

Pattern Classification

10. Linear Perceptron, Least Squares & Multi-layer NNs

AbdElMoniem Bayoumi, PhD

Fall 2021

Recap: Linear Perceptron Algorithm

1. Initialize the weights and threshold (bias) randomly
2. Present the augmented input (or feature) vector of the m^{th} training $\underline{u}(m)$ and its corresponding desired output $d(m)$

$$d(m) = \begin{cases} 1, & \text{if } \underline{u}(m) \in C_1 \\ 0, & \text{if } \underline{u}(m) \in C_2 \end{cases}$$

3. Calculate the actual output for pattern m :

$$y(m) = f(\underline{W}^T \underline{u}(m))$$

4. Adapt the weights according to the following rule (called Widrow-Hoff rule):

$$\underline{W}(\text{new}) = \underline{W}(\text{old}) + \alpha[d(m) - y(m)]\underline{u}(m)$$

where α is a constant called the learning rate

5. Go to step 2 until all patterns are classified correctly, i.e., $d(m)=y(m)$ for $m=1, \dots, M$

Note: the algorithm is sequential w.r.t. the patterns $\underline{u}(m)$

Recap: Understanding Widrow-Hoff Update

- If $d(m)=y(m)$ then no change is needed in the weights, i.e., $\underline{W}(new) = \underline{W}(old)$, because $d(m)-y(m)=0$
- If $d(m) \neq y(m)$ then weights get updated

$$\underline{W}(new) = \underline{W}(old) + \alpha[d(m) - y(m)]\underline{u}(m)$$

Linear Perceptron Algorithm

$$\underline{W}(new) = \underline{W}(old) + \alpha[d(m) - y(m)]\underline{u}(m)$$

- To show that each iteration corrects errors:
 - Let actual class $d(m) = 1$
 - If neuron classification $y(m) = f(\underline{W}^T \underline{u}(m)) = 0$
 - So, $\underline{W}^T \underline{u}(m) < 0$
 - However, we want $y(m) = 1$, i.e., $y(m) = d(m)$
 - We need to correct wrong classification by making what is inside $f(\cdot)$ more positive, which will make $y(m)$ move likely to be 1

Linear Perceptron Algorithm

$$\underline{W}(new) = \underline{W}(old) + \alpha[d(m) - y(m)]\underline{u}(m)$$

$$y(new) = f(\underline{W}^T(new) \cdot \underline{u}(m))$$

$$= f(\underline{W}^T(old) \cdot \underline{u}(m) + \alpha[d(m) - y(old)]\underline{u}^T(m)\underline{u}(m))$$

$$= f(\underline{W}^T(old) \cdot \underline{u}(m) + \alpha[1 - 0]\underline{u}^T(m)\underline{u}(m))$$

$$= f(-ve + \alpha\|\underline{u}(m)\|^2)$$

+ve

$\alpha > 0$

- Tries to make what is inside $f(\cdot)$ more positive, which will make $y(m)$ move likely to be 1

Theorem (Rosenblatt)

- For a linearly separable problem the perceptron algorithm is guaranteed to converge, leading to a solution that classifies all points correctly
- If the problem is not linearly separable, then the algorithm will not converge & will keep cycling forever
- How to deal with not linearly separable problems?

Least Square Classifier

- We try to have the neuron produce positive numbers for patterns from class 1 & negative numbers for patterns from class 2

$$\underline{X}(m) \rightarrow b_m$$

- $b_m > 0$ if $\underline{X}(m) \in C_1$
- $b_m < 0$ if $\underline{X}(m) \in C_2$
- $y(m) = f(\underline{W}^T \underline{u}(m)) = \underline{W}^T \underline{u}(m)$

linear activation fn.

Least Square Classifier

- Example:

- $\underline{X}(1) = \begin{bmatrix} 1 \\ -2 \\ 3 \end{bmatrix}, \quad \underline{X}(1) \in C_1, \quad b_1 = 1$
- $\underline{X}(2) = \begin{bmatrix} 3 \\ -2 \\ 5 \end{bmatrix}, \quad \underline{X}(2) \in C_1, \quad b_2 = 2$
- $\underline{X}(3) = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}, \quad \underline{X}(3) \in C_2, \quad b_3 = -3$
- \vdots

- We want:

- $\underline{W}^T \underline{u}(1) \approx b_1 \equiv 1$
- $\underline{W}^T \underline{u}(2) \approx b_2 \equiv 2$
- $\underline{W}^T \underline{u}(3) \approx b_3 \equiv -3$

Find W that satisfy these equations!

- If **M > N + 1**, then we cannot get exact equality because

#unknowns < #equations

$$\underline{W} \rightarrow N + 1 \quad M$$

Least Square Classifier

- Define error function:

$$E = \sum_{m=1}^M (\underline{W}^T \underline{u}(m) - b_m)^2$$

- It measures how close the obtained solution is to the desired one.
- We then seek to minimize the error fn.
- Thus, we try to find \underline{W} that minimizes E

Least Square Classifier

- Define error function:

$$E = \sum_{m=1}^M (\underline{W}^T \underline{u}(m) - b_m)^2$$


At the minimum:

$$\frac{\partial E}{\partial W_0} = 0$$

$$\frac{\partial E}{\partial W_1} = 0$$

\vdots

$$\frac{\partial E}{\partial W_N} = 0$$



Least Square Classifier

- Define the gradient vector:

$$\frac{\partial E}{\partial \underline{W}} = \begin{bmatrix} \frac{\partial E}{\partial W_0} \\ \frac{\partial E}{\partial W_1} \\ \vdots \\ \frac{\partial E}{\partial W_N} \end{bmatrix}$$

- Set $\frac{\partial E}{\partial \underline{W}} = \mathbf{0}$ and solve for \underline{W}
- **Advantage:**
 - Can converge if the problem is not linearly separable
- **Disadvantage:**
 - Linear classifier \rightarrow not suitable for most applications

Least Square Classifier

- $E = \sum_{m=1}^M (\underline{W}^T \underline{u}(m) - b_m)^2 = \sum_{m=1}^M \underbrace{(\underline{u}^T(m) \underline{W} - b_m)}_{\mathbf{Z}_m}^2$
- Let $Y = \begin{bmatrix} \underline{u}^T(1) \\ \underline{u}^T(2) \\ \vdots \\ \underline{u}^T(M) \end{bmatrix}$ **Matrix M x N**
- Then, $Y \underline{W} - \underline{b} = \begin{bmatrix} \underline{u}^T(1) \\ \underline{u}^T(2) \\ \vdots \\ \underline{u}^T(M) \end{bmatrix} \underline{W} - \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_M \end{bmatrix} = \begin{bmatrix} \underline{u}^T(1) \underline{W} - b_1 \\ \underline{u}^T(2) \underline{W} - b_2 \\ \vdots \\ \underline{u}^T(M) \underline{W} - b_M \end{bmatrix} \equiv \underline{Z}$
- $$\begin{aligned}
 E &= \sum_{m=1}^M Z_m^2 = \|\underline{Z}\|^2 = \underline{Z}^T \underline{Z} = [\underline{Y} \underline{W} - \underline{b}]^T [\underline{Y} \underline{W} - \underline{b}] \\
 &= [\underline{W}^T \underline{Y}^T - \underline{b}^T] [\underline{Y} \underline{W} - \underline{b}] \\
 &= \underline{W}^T \underline{Y}^T \underline{Y} \underline{W} - \underline{W}^T \underline{Y}^T \underline{b} - \underline{b}^T \underline{Y} \underline{W} + \underline{b}^T \underline{b} \\
 &= \underline{W}^T \underline{Y}^T \underline{Y} \underline{W} - 2 \underline{b}^T \underline{Y} \underline{W} + \underline{b}^T \underline{b}
 \end{aligned}$$

Least Square Classifier

- $E = \underline{W}^T Y^T Y \underline{W} - 2 \underline{b}^T Y \underline{W} + \underline{b}^T \underline{b}$
- $\frac{\partial E}{\partial \underline{W}} = \mathbf{0}$
- $\frac{\partial}{\partial \underline{W}} [\underbrace{\underline{W}^T Y^T Y \underline{W}}_A - \underbrace{2 \underline{b}^T Y \underline{W}}_{C^T} + \underline{b}^T \underline{b}] = \mathbf{0}$
- $\underline{W}^T A \underline{W} = W_1^2 A_{11} + W_1 W_2 A_{12} + \dots$
 $\quad \quad \quad + W_1 W_2 A_{21} + W_2^2 A_{22} + \dots$
 $\quad \quad \quad + \dots$
- $\frac{\partial}{\partial \underline{W}} [\underline{W}^T A \underline{W}] = 2A\underline{W}$ $\frac{\partial}{\partial \underline{W}} [C^T \underline{W}] = C$

Quadratic form

Exercise!

Least Square Classifier

- $$\begin{aligned}\frac{\partial E}{\partial \underline{W}} &= \frac{\partial}{\partial \underline{W}} [\underline{W}^T Y^T Y \underline{W} - 2 \underline{b}^T Y \underline{W} + \underline{b}^T \underline{b}] \\ &= 2A\underline{W} - \underline{C} + \underline{0} \\ &= 2Y^T Y \underline{W} - 2Y^T b = \underline{0}\end{aligned}$$
- $\underline{W} = (Y^T Y)^{-1} Y^T b$ **Weights for the least square classifier**
- Advantage:
 - Can converge if the problem is not linearly separable


Multi-Layer Networks

- We have seen that a single neuron, or a single layer network is capable of only producing linear classifiers
- Hence, it is not adequate for most applications
- Some very simple fn.'s like XOR fn. Cannot be implemented with a single neuron


Multi-Layer Networks

- XOR function:

Input 1	Input 2	Output
0	0	0
0	1	1
1	0	1
1	1	0




Multi-Layer Networks




- We need non-linear decision region \rightarrow use multi-layer network

Multi-Layer Networks




Multi-Layer Networks




Example on Multi-Layer Networks

- A one-hidden-layer network can implement the XOR problem
- Note: We need two linear classifiers to solve the problem



Example on Multi-Layer Networks



$$Z_1 = f(w_{11}X_1 + w_{21}X_2 + w_{10}) = f(2X_1 - 2X_2 - 1)$$

$$Z_2 = f(w_{12}X_1 + w_{22}X_2 + w_{20}) = f(-2X_1 + 2X_2 - 1)$$

$$Y = f(v_1Z_1 + v_2Z_2 + v_0) = f(Z_1 + Z_2 - 0.5)$$

Hidden neurons output

Network output

$f(.)$ is a step function

Example on Multi-Layer Networks

$$Z_1 = f(w_{11}X_1 + w_{12}X_2 + w_{10}) = f(2X_1 - 2X_2 - 1)$$

$$Z_2 = f(w_{21}X_1 + w_{22}X_2 + w_{20}) = f(-2X_1 + 2X_2 - 1)$$

$$Y = f(v_1Z_1 + v_2Z_2 + v_0) = f(Z_1 + Z_2 - 0.5)$$

X1	X2	Z1	Z2	Y	d target
0	0	0	0	0	0
0	1	0	1	1	1
1	0	1	0	1	1
1	1	0	0	0	0


This network can implement the XOR fn. correctly

Training Multi-Layer Networks


- We have seen that multi-layer networks can implement non-linear fn.'s /decision regions
- The powerful feature of multilayer networks (aka. feed-forward networks) is its ability to learn
- Again, we use a training set collected from the problem we wish to solve, i.e., $\underline{u}(m)$ & $d(m)$

Training Multi-Layer Networks

- The target output $d(m)$ could be the classification of a pattern in case of pattern classification problem
- Alternatively, the target output could be the actual value to be predicted by the NN




Training Multi-Layer Networks



- We usually use hidden node fn's that are continuous (like the ReLu fn.)


Activation Functions

- ReLU (Rectified Linear Unit)



$$f(x) = \max(x, 0)$$

- Leaky ReLU




$$f(x) = \begin{cases} x & \text{if } x > 0 \\ ax & \text{otherwise} \end{cases}$$

where $0 < a < 1$


Activation Functions

- Tanh



$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

- Sigmoid (used only in o/p layer in binary classification problems)



$$f(x) = \frac{1}{1 + e^{-x}}$$

Activation Functions

- ReLU & Leaky ReLU are most famous now
 - Especially with learning from images
 - Faster learning
- Tanh
 - comes next after ReLU
 - More famous with sequence problems, e.g., speech recognition
- Sigmoid
 - Slow learning
 - Used only in output layer in binary classification problems

Supervised Learning

- Learning in case of input/output training examples
- We would like the network to produce an output $y(m)$ when inputting $u(m)$ as close as possible to the target output $d(m)$
- We define an error function, i.e., cost function, as a measure of how close the network outputs to the target outputs
 - Choice depends on the problem
 - E.g. MSE (Mean Square Error):

$$E = \frac{1}{M} \sum_{m=1}^M (y(m) - d(m))^2$$

Supervised Learning

- How learning is done is that we adjust the weights in small steps, so that each step decreases the error function a little
- We keep repeating this process until the error reaches its lowest value, i.e., at which $y(m)$ will be as close as possible to $d(m)$
- Our goal is to minimize $E \rightarrow$ find weights that give minimum E

Supervised Learning

- Denote $\underline{W} = \begin{bmatrix} w_{11} \\ w_{12} \\ \vdots \end{bmatrix} \rightarrow \text{get min } E(\underline{W})$
- At the minimum $\frac{\partial E}{\partial w_{11}} = 0, \frac{\partial E}{\partial w_{12}} = 0 \dots$

$$-\frac{\partial E}{\partial \underline{W}} = \begin{bmatrix} \frac{\partial E}{\partial w_{11}} \\ \frac{\partial E}{\partial w_{12}} \\ \vdots \end{bmatrix} \quad \text{gradient vector}$$

- $\frac{\partial E}{\partial \underline{W}} = 0 \rightarrow$ has no analytical or closed form solution, i.e., cannot be algebraically solved

Gradient Descent



- We use the concept of steepest descent, aka. gradient descent
- It is a general algorithm for minimizing fn's
- We update the weights as:

$$\underline{W}(new) = \underline{W}(old) - \alpha \frac{\partial E}{\partial \underline{W}}$$

- α is a constant, i.e., learning rate that determines the size of the move

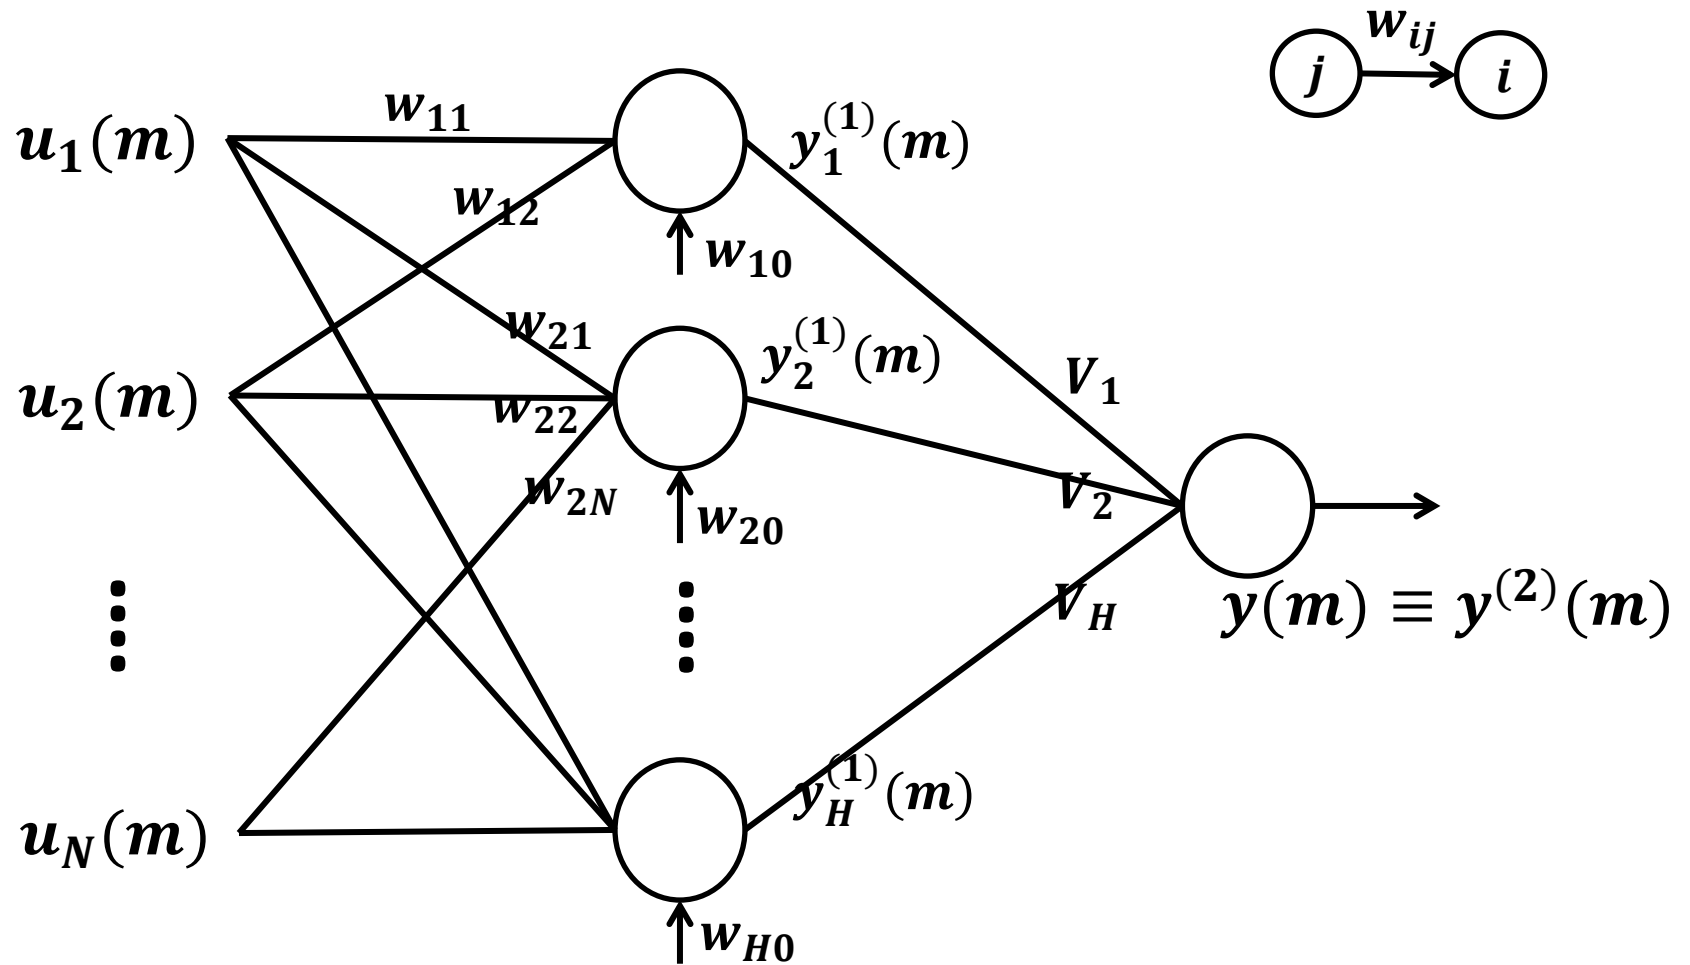
Gradient Descent

- It can be shown that the negative direction of the gradient gives the steepest descent



- When we approach the min, the steps become very small because close to the min we find $\frac{\partial E}{\partial W} \approx 0$

Example



Let $\underline{y}^{(0)}(m) \equiv \underline{u}(m)$

$H \equiv$ no. of hidden nodes

Example

- $E = \sum_{m=1}^M [y^{(2)}(m) - d(m)]^2 = \sum_{m=1}^M e^2(m)$
- To get gradient, compute:
$$-\frac{\partial E}{\partial W_{IJ}} \quad \& \quad \frac{\partial E}{\partial V_I}$$
- $y_i^{(1)}(m) = f\left(\sum_{j=1}^N w_{ij} u_j(m) + w_{i0}\right)$
- $y^{(2)}(m) = f\left(\sum_{j=1}^H V_j y_j^{(1)}(m) + V_0\right)$

Example

- $E = \sum_{m=1}^M [y^{(2)}(m) - d(m)]^2 = \sum_{m=1}^M e^2(m)$
- $\frac{\partial E}{\partial V_I} = \sum_{m=1}^M 2 e(m) \frac{\partial e(m)}{\partial V_I}$
- $$\begin{aligned} \frac{\partial e(m)}{\partial V_I} &= \frac{\partial [y^{(2)}(m) - d(m)]}{\partial V_I} = \frac{\partial y^{(2)}(m)}{\partial V_I} \\ &= f' \left(\sum_{j=1}^H V_j y_j^{(1)}(m) + V_0 \right) \frac{\partial [V_1 y_1^{(1)}(m) + \dots + V_H y_H^{(1)}(m) + V_0]}{\partial V_I} \\ &= f' \left(\sum_{j=1}^H V_j y_j^{(1)}(m) + V_0 \right) \frac{\partial V_I y_I^{(1)}(m)}{\partial V_I} = f' \left(\sum_{j=1}^H V_j y_j^{(1)}(m) + V_0 \right) y_I^{(1)}(m) \end{aligned}$$
- $$\begin{aligned} \frac{\partial E}{\partial V_I} &= \sum_{m=1}^M 2 e(m) f' \left(\sum_{j=1}^H V_j y_j^{(1)}(m) + V_0 \right) y_I^{(1)}(m) \\ &= \sum_{m=1}^M 2 e(m) f' \left(\sum_{j=1}^H V_j y_j^{(1)}(m) + V_0 \right) f \left(\sum_{j=1}^N w_{Ij} u_j(m) + w_{I0} \right) \end{aligned}$$

Example

- $E = \sum_{m=1}^M [y^{(2)}(m) - d(m)]^2 = \sum_{m=1}^M e^2(m)$
- $\frac{\partial E}{\partial w_{IJ}} = \sum_{m=1}^M 2 e(m) \frac{\partial e(m)}{\partial w_{IJ}}$
- $$\begin{aligned} \frac{\partial e(m)}{\partial w_{IJ}} &= \frac{\partial [y^{(2)}(m) - d(m)]}{\partial w_{IJ}} = \frac{\partial y^{(2)}(m)}{\partial w_{IJ}} \\ &= f' \left(\sum_{j=1}^H V_j y_j^{(1)}(m) + V_0 \right) \frac{\partial [V_1 y_1^{(1)}(m) + \dots + V_H y_H^{(1)}(m) + V_0]}{\partial w_{IJ}} \\ &= f' \left(\sum_{j=1}^H V_j y_j^{(1)}(m) + V_0 \right) \frac{\partial [V_1 f(\sum_{j=1}^N w_{1j} u_j(m) + w_{10}) + \dots + V_H f(\sum_{j=1}^N w_{Hj} u_j(m) + w_{H0}) + V_0]}{\partial w_{IJ}} \\ &= f' \left(\sum_{j=1}^H V_j y_j^{(1)}(m) + V_0 \right) \frac{\partial V_I f(\sum_{j=1}^N w_{Ij} u_j(m) + w_{I0})}{\partial w_{IJ}} \\ &= f' \left(\sum_{j=1}^H V_j y_j^{(1)}(m) + V_0 \right) V_I f' \left(\sum_{j=1}^N w_{Ij} u_j(m) + w_{I0} \right) \frac{\partial [\sum_{j=1}^N w_{Ij} u_j(m) + w_{I0}]}{\partial w_{IJ}} \\ &= f' \left(\sum_{j=1}^H V_j y_j^{(1)}(m) + V_0 \right) V_I f' \left(\sum_{j=1}^N w_{Ij} u_j(m) + w_{I0} \right) u_j(m) \end{aligned}$$
- $$\frac{\partial E}{\partial w_{IJ}} = \sum_{m=1}^M 2 e(m) f' \left(\sum_{j=1}^H V_j y_j^{(1)}(m) + V_0 \right) V_I f' \left(\sum_{j=1}^N w_{Ij} u_j(m) + w_{I0} \right) u_j(m)$$

Acknowledgment

- These slides have been created relying on lecture notes of Prof. Dr. Amir Atiya