

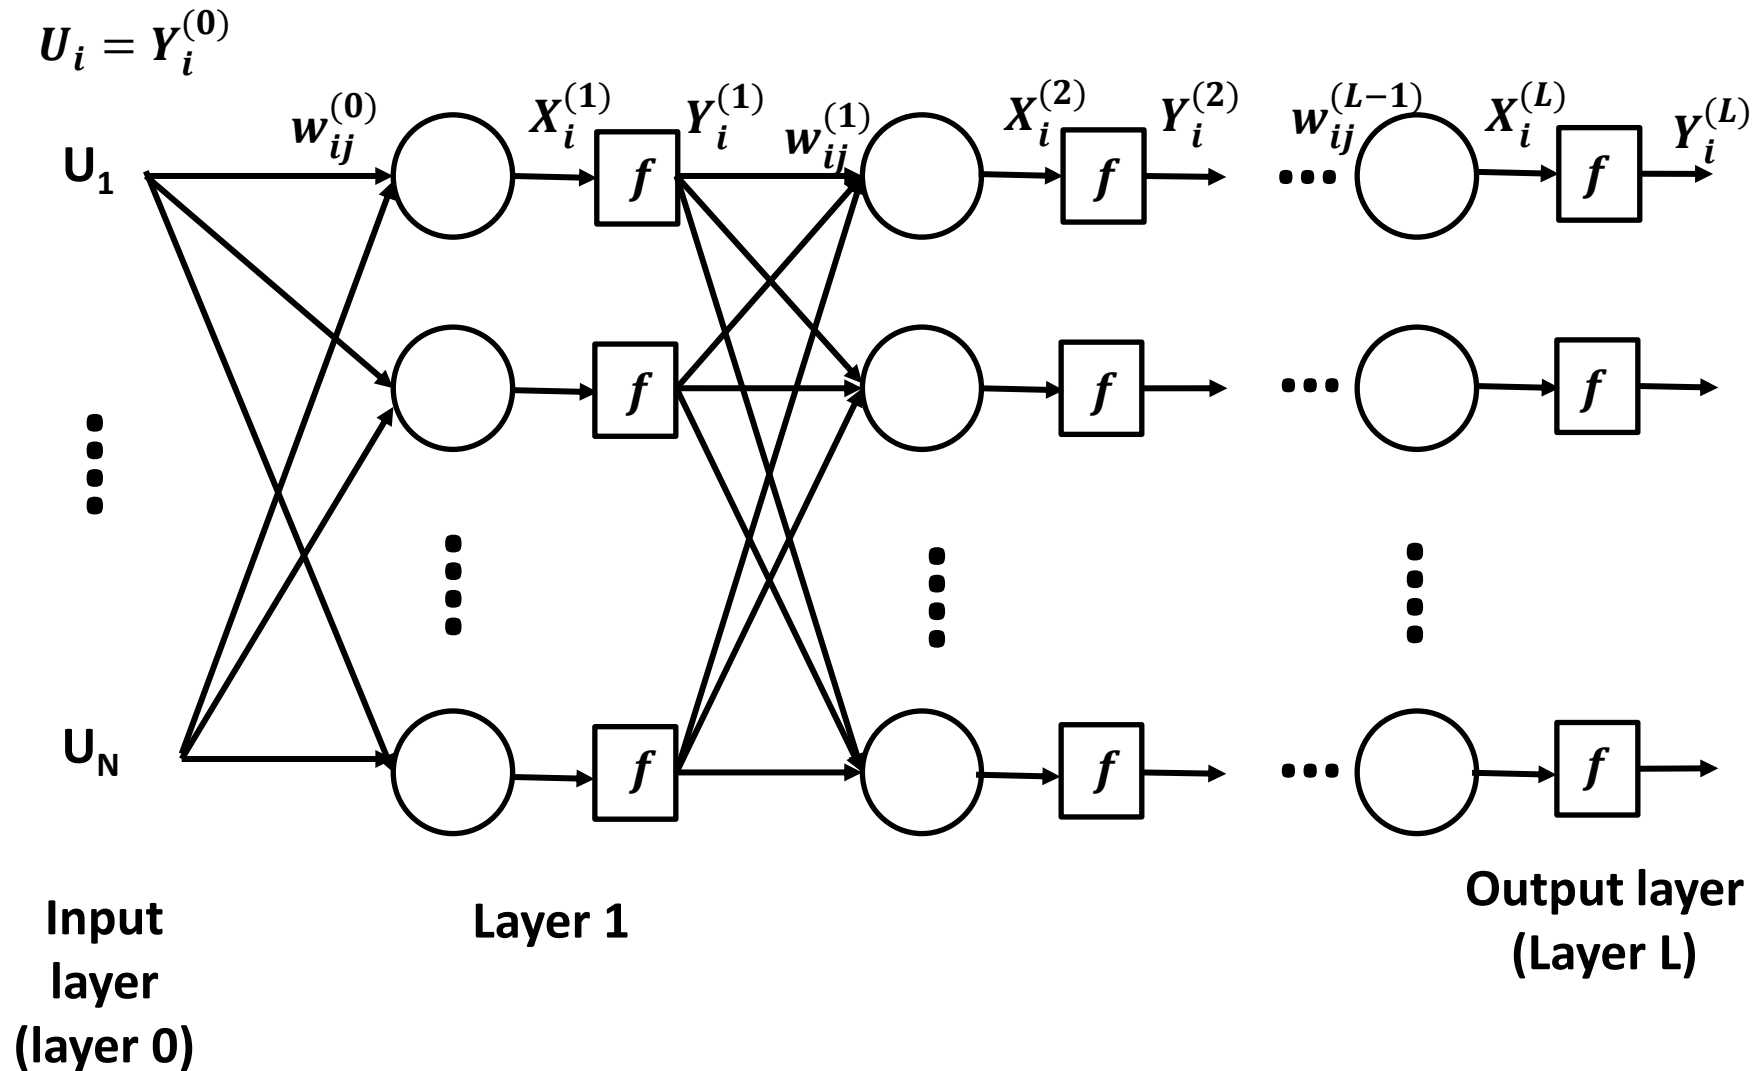
Pattern Classification

11. Backpropagation & Time-Series Forecasting

AbdElMoniem Bayoumi, PhD

Spring 2022



Recap: Multi-Layer Networks



$N(l)$ is the number of nodes in layer l

Recap: Gradient Descent

- It can be shown that the negative direction of the gradient gives the steepest descent



- When we approach the min, the steps become very small because close to the min we find $\frac{\partial E}{\partial W} \approx 0$

Back propagation Algorithm

- It is an algorithm based on the steepest descent concept
- Used to train a general multi-layer network

Back propagation Algorithm

- $X_i^{(l)} = \sum_{j=1}^{N^{(l-1)}} w_{ij}^{(l-1)} Y_j^{(l-1)}$ **Output of hidden node before applying the activation function**

- $Y_i^{(l)} = f\left(X_i^{(l)}\right)$ **Output of layer**


Back propagation Algorithm

Error, i.e., cost fn.

- $E = \frac{1}{M} \sum_{m=1}^M E_m$
- $E_m = \sum_{i=1}^{N(L)} \left[Y_i^{(L)}(m) - d_i(m) \right]^2$ Loss (i.e., for regression)
- $Y_i^{(L)}(m) \equiv i^{th}$ output of the NN for the training pattern m
- $d_i(m) \equiv$ target o/p
- $N(L) \equiv$ no. of outputs (i.e., nodes of the output layer)
- We need to compute the gradient, i.e., $\frac{\partial E}{\partial w_{ij}^{(l)}}$

Chain Rule

- $Y = f(y_1, y_2, y_3)$
- $y_1 = g_1(Z)$, $y_2 = g_2(Z)$, $y_3 = g_3(Z)$



- $$\frac{\partial Y}{\partial Z} = \frac{\partial Y}{\partial y_1} * \frac{\partial y_1}{\partial Z} + \frac{\partial Y}{\partial y_2} * \frac{\partial y_2}{\partial Z} + \frac{\partial Y}{\partial y_3} * \frac{\partial y_3}{\partial Z}$$

Back propagation Algorithm

- $E_m = \sum_{i=1}^{N^{(L)}} \left[Y_i^{(L)}(m) - d_i(m) \right]^2$
- $Y_i^{(L)} = f \left(X_i^{(L)} \right)$
- $X_i^{(L)} = \sum_{j=1}^{N^{(L-1)}} w_{ij}^{(L-1)} Y_j^{(L-1)}$


$$\begin{aligned} \frac{\partial E_m}{\partial w_{IJ}^{(L-1)}} &= \frac{\partial E_m}{\partial Y_I^{(L)}} * \frac{\partial Y_I^{(L)}}{\partial X_I^{(L)}} * \frac{\partial X_I^{(L)}}{\partial w_{IJ}^{(L-1)}} \\ &= \frac{\partial E_m}{\partial Y_I^{(L)}} * \frac{\partial Y_I^{(L)}}{\partial X_I^{(L)}} * \frac{\partial \left(\sum_{j=1}^{N^{(L-1)}} w_{IJ}^{(L-1)} Y_j^{(L-1)} \right)}{\partial w_{IJ}^{(L-1)}} \\ &= \frac{\partial E_m}{\partial Y_I^{(L)}} * \frac{\partial Y_I^{(L)}}{\partial X_I^{(L)}} * Y_J^{(L-1)} \end{aligned}$$

Chain rule!

- Note that the other $X_i^{(L)}$'s are not taken into account, because they do not depend on $w_{IJ}^{(L-1)}$ at all

Back propagation Algorithm

- To get $\frac{\partial E_m}{\partial w_{ij}^{(l)}}$ for any general layer l




$$\frac{\partial E_m}{\partial w_{IJ}^{(l)}} = \frac{\partial E_m}{\partial X_I^{(l+1)}} * \frac{\partial X_I^{(l+1)}}{\partial w_{IJ}^{(l)}} = \frac{\partial E_m}{\partial X_I^{(l+1)}} * Y_J^{(l)}$$

$$\frac{\partial E_m}{\partial X_I^{(l+1)}} = \sum_{i=1}^{N^{(l+2)}} \frac{\partial E_m}{\partial X_i^{(l+2)}} * \frac{\partial X_i^{(l+2)}}{\partial X_I^{(l+1)}} = \sum_{i=1}^{N^{(l+2)}} \frac{\partial E_m}{\partial X_i^{(l+2)}} * \frac{\partial X_i^{(l+2)}}{\partial Y_I^{(l+1)}} * \frac{\partial Y_I^{(l+1)}}{\partial X_I^{(l+1)}}$$

Back propagation Algorithm

1. Initialize all weights to small randomly chosen values, e.g. $[-1,1]$
2. Let $u(m)$ & $d(m)$ be the training input/output examples
3. For $m=1$ to M :
 - i. Present $u(m)$ to the network and compute the hidden layer outputs and final layer outputs
 - ii. Use these outputs in a backward scheme to compute the partial derivatives of error fn. w.r.t. to the weights of each layer
 - iii. Update weights: $w_{ij}^{[l]}(new) = w_{ij}^{[l]}(old) - \alpha \frac{\partial E_m}{\partial w_{ij}^{[l]}}$
4. Compute total error (stop in case of convergence)



Note: l refers to layer l

Disadvantages of Back propagation

- Can often be slow in reaching the min (i.e., sometimes tens of thousands of iterations)
 - Especially close to min
 - Too small $\alpha \rightarrow$ very small steps & slow to reach min
 - Too large $\alpha \rightarrow$ leads to oscillations & possibly not converging at all
 - Use variable α (start large then decrease it)

Disadvantages of Back propagation

- Prone to get stuck in local minima
 - This problem could be alleviated to some extent by repeating the training many times, each time from a different set of initial weights

Types of Weight Update

- Batch or epoch update, i.e, **Gradient Descent**:
 - Present **full set of training examples** (batch of examples)
 - Compute the error of each example
 - Compute gradient of the batch (based on cost function of the whole batch)
 - Update weights based on this batch gradient
 - Do another iteration ... and so on
 - Advantages:
 - Optimization is more consistent
 - Disadvantages:
 - Slow (too long per iteration)

Types of Weight Update


- Sequential update, i.e., **Stochastic Gradient Descent**:
 - Present a training pattern, then update the weights (according to $\frac{\partial E_m}{\partial w}$), then present the next one ... and so on
 - After finishing all patterns, do another iteration starting from $m = 1$
 - Advantages:
 - Faster compared to gradient descent, i.e., full-batch
 - Disadvantages:
 - Hard to converge: “stochastic” since it depends on every single example; however, **in practice** being close to minimum is **reasonably good**
 - Loss speedup from vectorization
 - In practice for large datasets SGD is preferred to GD

Types of Weight Update

- Mini-Batch :
 - Present **subset** of training examples (mini-batch of examples)
 - Compute the error of each example
 - Compute gradient of the mini-batch (based on cost function of this mini-batch)
 - Update weights based on this mini-batch gradient
 - Move to another mini-batch & after finishing all mini-batches do another iteration ... and so on
 - Advantages:
 - Fast

Other Optimization Algorithms

- Gradient descent with momentum
 - Smooth-out the steps of the gradient descent using a moving average of the derivatives
 - Get faster learning in the intended direction & avoid oscillations



$$DW = \beta DW + (1 - \beta) \frac{\partial E}{\partial W}$$


$$W = W - \alpha DW$$

Faster learning
←→
Slower learning
↑↓

$DW = 0$ initially

Other Optimization Algorithms

- RMSProp
 - Slow-down learning in unintended directions
 - Avoid oscillations



$$Sw_i = \beta Sw_i + (1 - \beta) \left[\frac{\partial E}{\partial w_i} \right]^2$$

small

$$Sw_j = \beta Sw_j + (1 - \beta) \left[\frac{\partial E}{\partial w_j} \right]^2$$

large

$$\frac{\partial E}{\partial w_i} < \frac{\partial E}{\partial w_j}$$

$$w_i = w_i - \alpha \frac{\frac{\partial E}{\partial w_i}}{\sqrt{Sw_i} + \epsilon}$$

$$w_j = w_j - \alpha \frac{\frac{\partial E}{\partial w_j}}{\sqrt{Sw_j} + \epsilon}$$

Other Optimization Algorithms

- Adam (combines both RMSProp & momentum)

$$Dw_i = \beta_1 Dw_i + (1 - \beta_1) \frac{\partial E}{\partial w_i} \qquad Sw_i = \beta_2 Sw_i + (1 - \beta_2) \left[\frac{\partial E}{\partial w_i} \right]^2$$

$$w_i = w_i - \alpha \frac{Dw_i}{\sqrt{Sw_i} + \epsilon}$$

Regularization

- **Used to prevent overfitting**

- Intuition: set the weights of some hidden nodes to zero to simplify the network, i.e., smaller network

- L₂ regularization (aka weight decay): $J = \frac{1}{M} \sum_{m=1}^M E_m + \frac{\lambda}{2M} \|\underline{W}\|_2^2$

- $\|\underline{W}\|_2^2 = \sum_j w_j^2 = \underline{W}^T \underline{W}$

- L₁ regularization: $J = \frac{1}{M} \sum_{m=1}^M E_m + \frac{\lambda}{2M} \|\underline{W}\|_1$

- $\|\underline{W}\|_1 = \sum_j |w_j|$

- \underline{W} is the weights vector, thresholds not necessary to be included
- L₂ regularization is used more often
- λ is the regularization parameter (hyper-parameter to be tuned)

Dropout Regularization


- **Used to prevent overfitting**
- Intuitions:
 - Eliminate some nodes to simplify the network based on some probability, i.e., smaller network
 - As if you train smaller networks on individual training examples
 - Cannot rely on any one feature, so spread weights
- For each layer set a dropout probability
 - Each node within that layer may get eliminated based on that probability

Guidelines for Training

- Learning rate α :
 - Too small. Convergence will be slow.
 - Too large: we will oscillate around the minimum.
 - Some methods propose varying rate. i.e., learning rate decay.
 - When learning does not go well, consider using smaller learning rate.

Learning Rate Decay


- Gradient descent with small mini-batch size



Source: Andrew Ng

Input and Output Normalization

- Input and Output normalization
 - Inputs have to be approximately in the range of 0 to 1 or -1 to 1



- $x = (u - u_{\min}) / (u_{\max} - u_{\min})$
- $x = (u - \text{Mean}(u)) / \text{st dev}(u)$

Train/Dev/Test Partition

- Best practice:
 - **Training: 60% , Validation (Dev): 20% & Test: 20%**
 - In case of big data, e.g., 10^6 , then 98%, 1% & 1%
- Test set should be used only once, at the very end of the design

Machine Learning Recipe


- Train the network and evaluate first on the training data
 - If bias is high, i.e., underfitting (performance is bad on the training set itself), then:
 - **Bigger network** (more hidden nodes or more hidden layers) → works most of the time
 - Train longer → works sometimes
 - Check for bias again and keep changing until a good bias is reached
- Check for variance, i.e., performance on Dev set
 - If variance is high, i.e., overfitting (performance is bad on the validation set), then:
 - **More data** (if possible)
 - **Regularization**
 - Check again for bias first, then after that check for variance and so on until you reach a good bias & good variance
- Search for better NN architecture that better suits the problem (sometimes may work)

Hyper-Parameters Tuning

- Learning rate α **1st in importance**
- Momentum parameter $\beta \approx 0.9$
• Number of hidden nodes
• Mini-batch size **2nd in importance**
- Num of layers
• Learning-rate decay **3rd in importance**
- Adam parameters $\beta_1 \approx 0.9, \beta_2 \approx 0.999, \epsilon \approx 10^{-8}$
Not likely to make change!

Tuning Process


- Try random values: don't use a grid
 - Better exploration of important parameters
 - Consider the example on the board
- Coarse to fine scheme
 - Focus more on good regions
- Use appropriate scale
 - Do not sample uniformly
 - Use logarithmic scale
 - E.g., α range is $[0.0001, 1]$ linear scale scaling will give more weight to the values between 0.1 & 1, however, logarithmic scale:



Sample $r \sim [-4, 0]$
then $\alpha = 10^r$

Tuning Process

- Use appropriate scale
 - More example: let β range is $[0.9, 0.999]$
 - Sample from $1 - \beta$, i.e., $[0.1, 0.001]$, using log scale




- Sensitivity of β approaching has huge impact on the performance, i.e., momentum corresponds to averaging over the last $\frac{1}{1-\beta}$ examples
 - $\beta \sim [0.900, 0.9005] \rightarrow$ averaging over last 10 examples
 - $\beta \sim [0.999, 0.9995] \rightarrow$ 1000 to 2000 examples

K-Fold Cross Validation

- For parameter tuning over the training data
- Apply K-fold validation to the training set (usually $k=5$).

- Example $K=4$

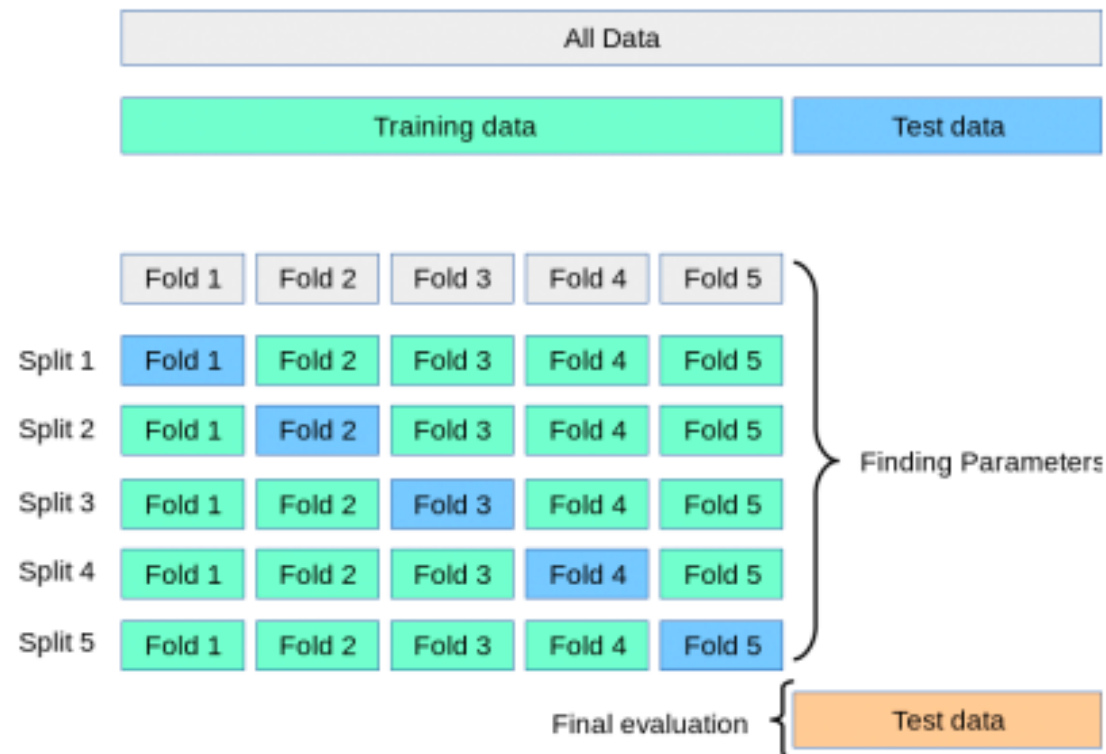
$$E_{VAL} = E_1 + E_2 + E_3 + E_4$$



- Repeat for every parameter value, minimize E_{VAL}


K-Fold Cross Validation

- Better than convention train-validation-test split
 - Just split not training and test (no need for validation set)
 - Not biased to the nature of splitting of the training and validation




Multi-Class Classification

- E.g., an image is either of a cat, or dog, or duck or otherwise




Multi-Class Classification

- Use sigmoid activation function in the output only in case of binary classification, i.e., two classes
- For multi-class classification use soft-max regression:



Multi-Class Classification

- For multi-class classification use soft-max regression:



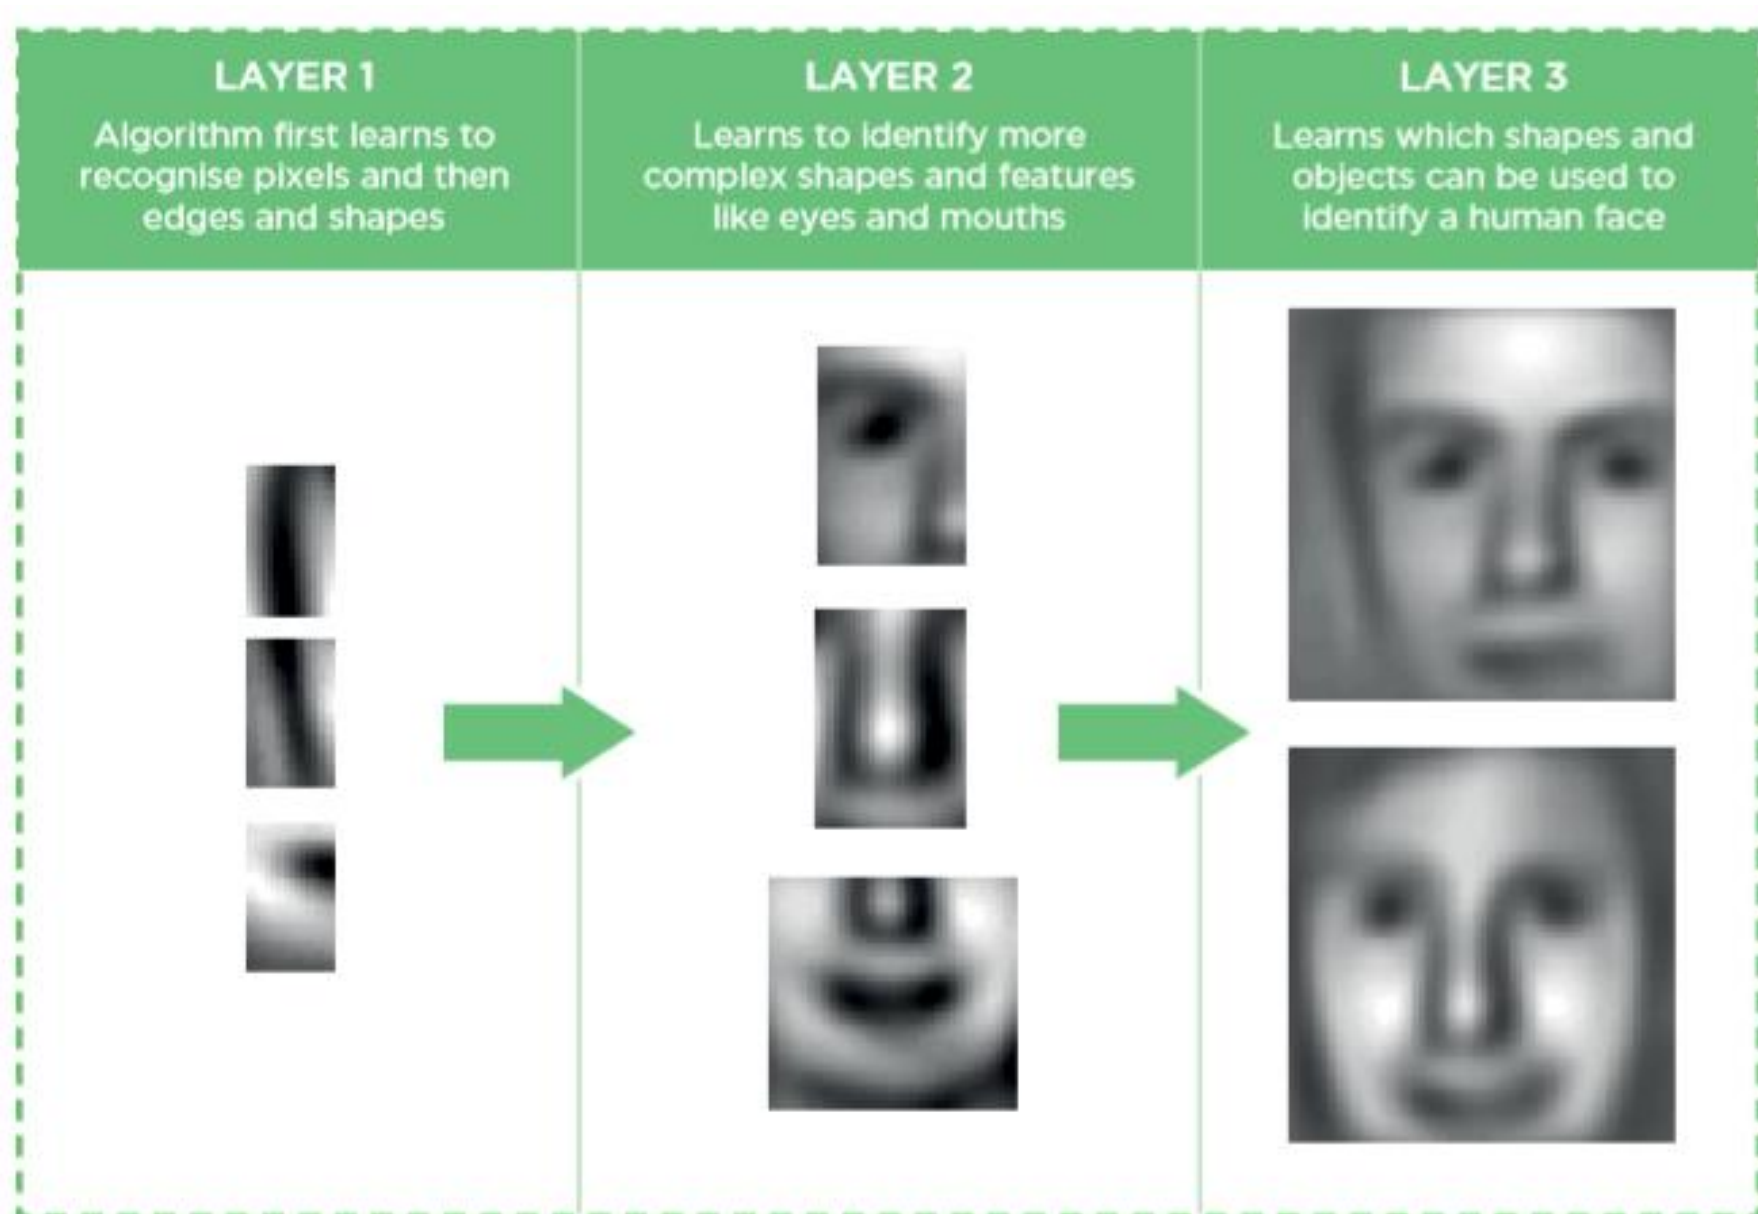
- $z_i^{[L]}$ output of node i at the output layer before applying any activation function

- y_i is the output after applying soft-max
- $$y_i = \frac{e^{z_i^{[L]}}}{\sum_j e^{z_j^{[L]}}}$$

Convolutional Neural Networks (CNN)

- Mostly applied to imagery problems
- Layers extract features from input images, e.g., edge detection

Convolutional Neural Networks (CNN)



Convolutional Neural Networks (CNN)

- Mostly applied to imagery problems
- Layers extract features from input images, e.g., edge detection
 - Convolution layer, i.e., filtering
 - Pooling Layer, i.e., reduce input (avg or max)
 - Fully Connected Layer, i.e., as in multi-layer NN, at the final layers

Vertical Edge Detection

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

*

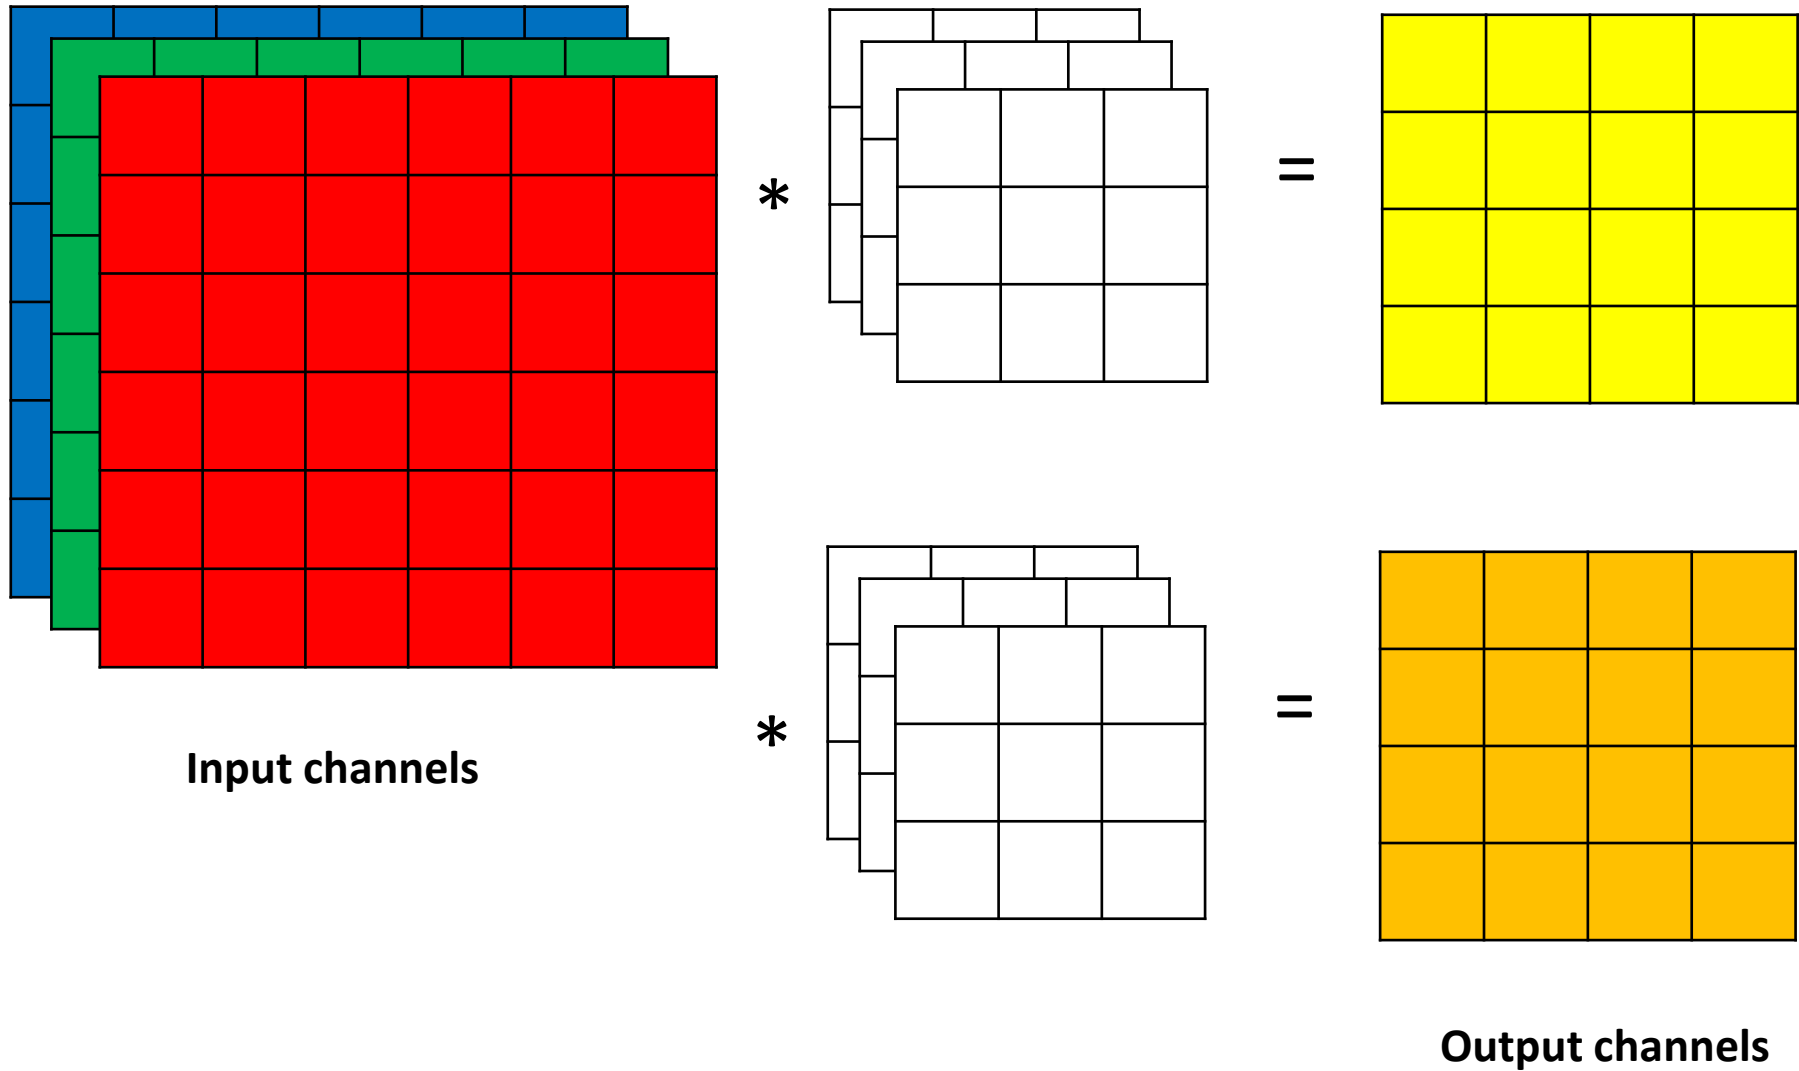
1	0	-1
1	0	-1
1	0	-1

=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0

convolution

Convolutional Neural Networks (CNN)



Max Pooling


10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0



10	10	0
10	10	0
10	10	0

What about average pooling?


Global Max Pooling



What about global average pooling?

Convolutional Neural Networks (CNN)

- Learn filters' parameters and weights of fully connected layers



Source: Andrew Ng's Lectures

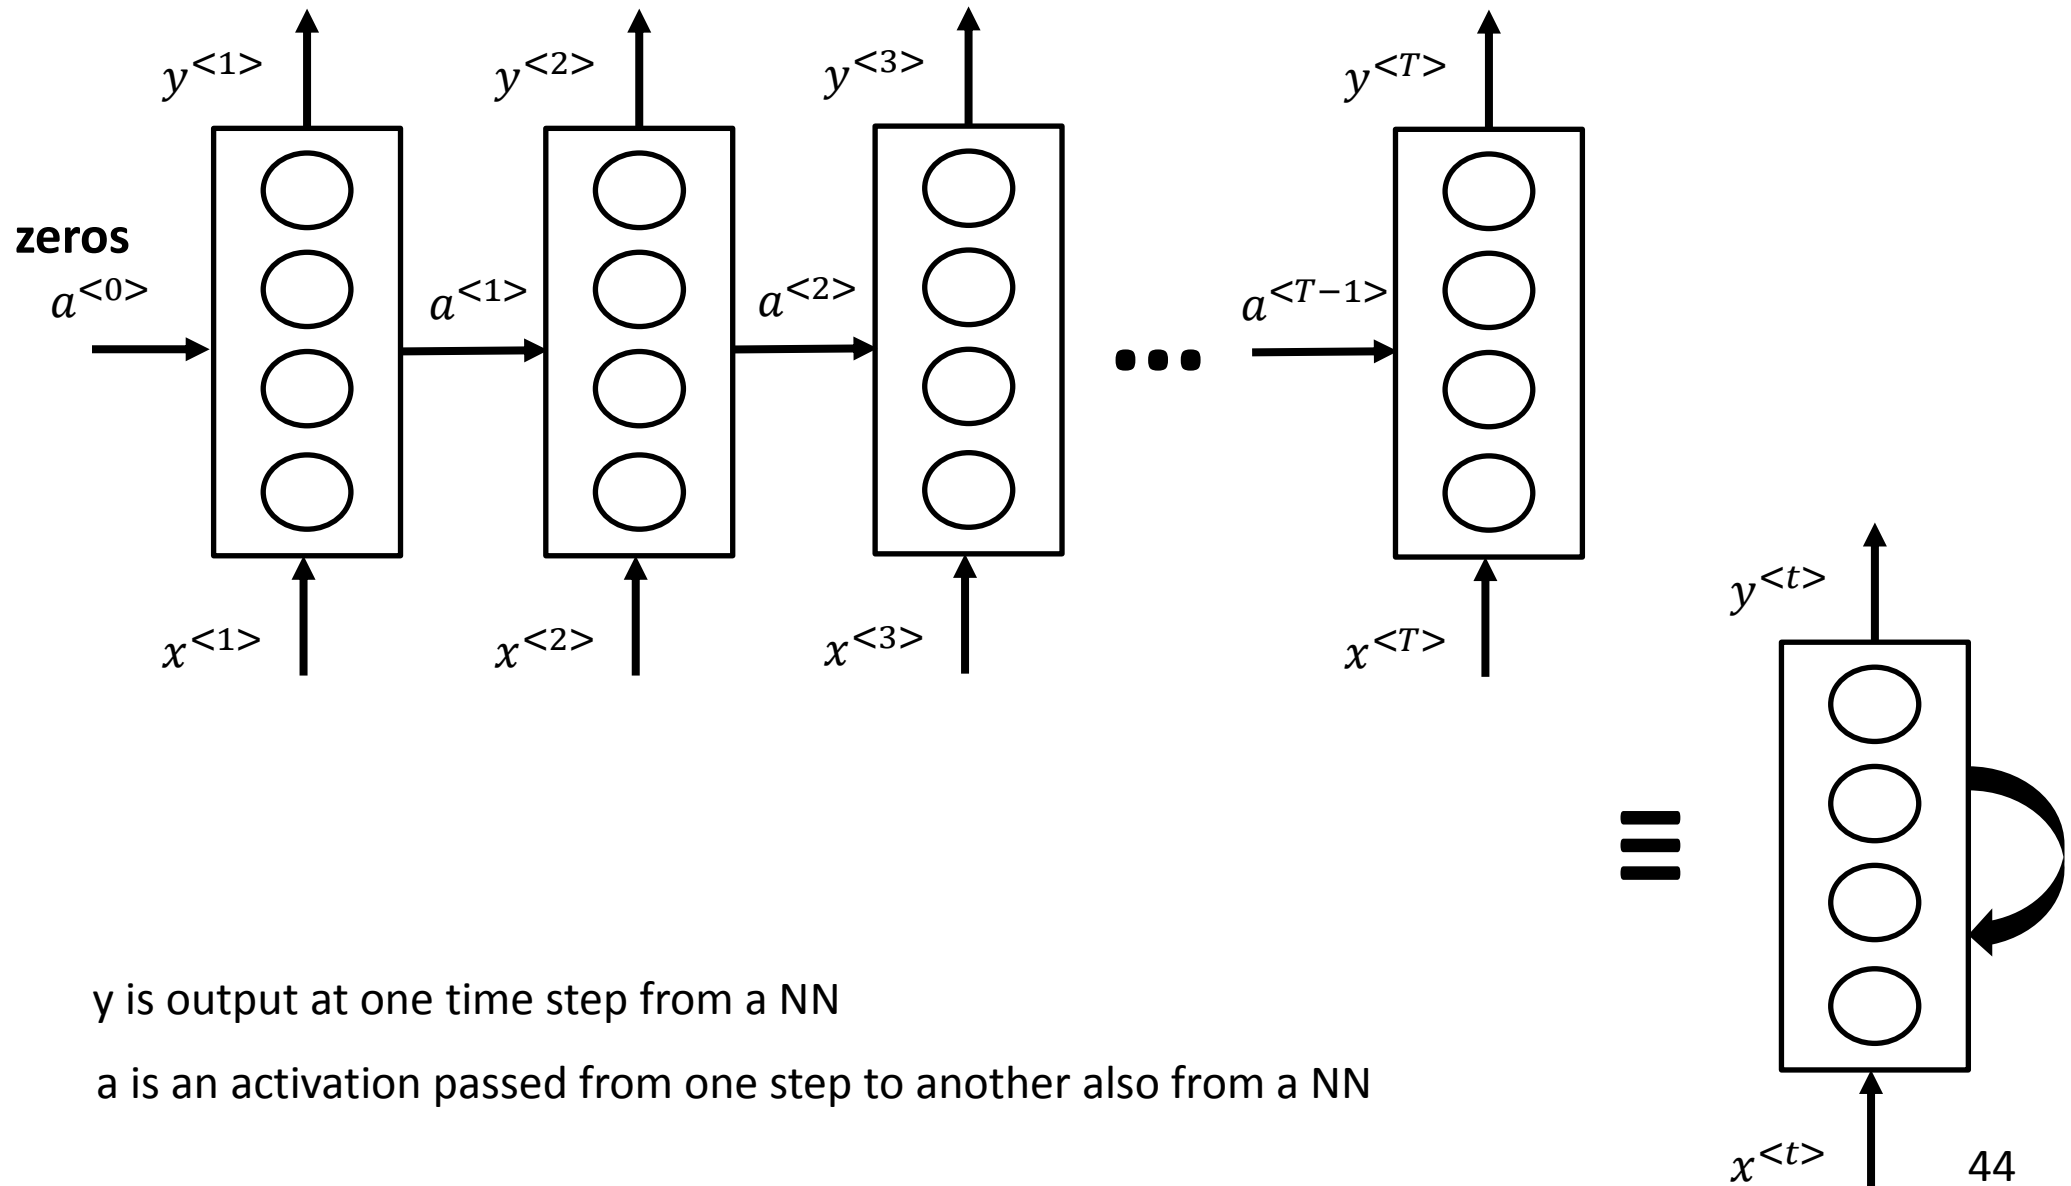
Convolutional Neural Networks (CNN)

- Convolution leads to less memory footprint due to:
 - Parameter sharing (compared to fully connected layers)
 - Sparsity of connections (at each layer output depends on limited number of inputs)

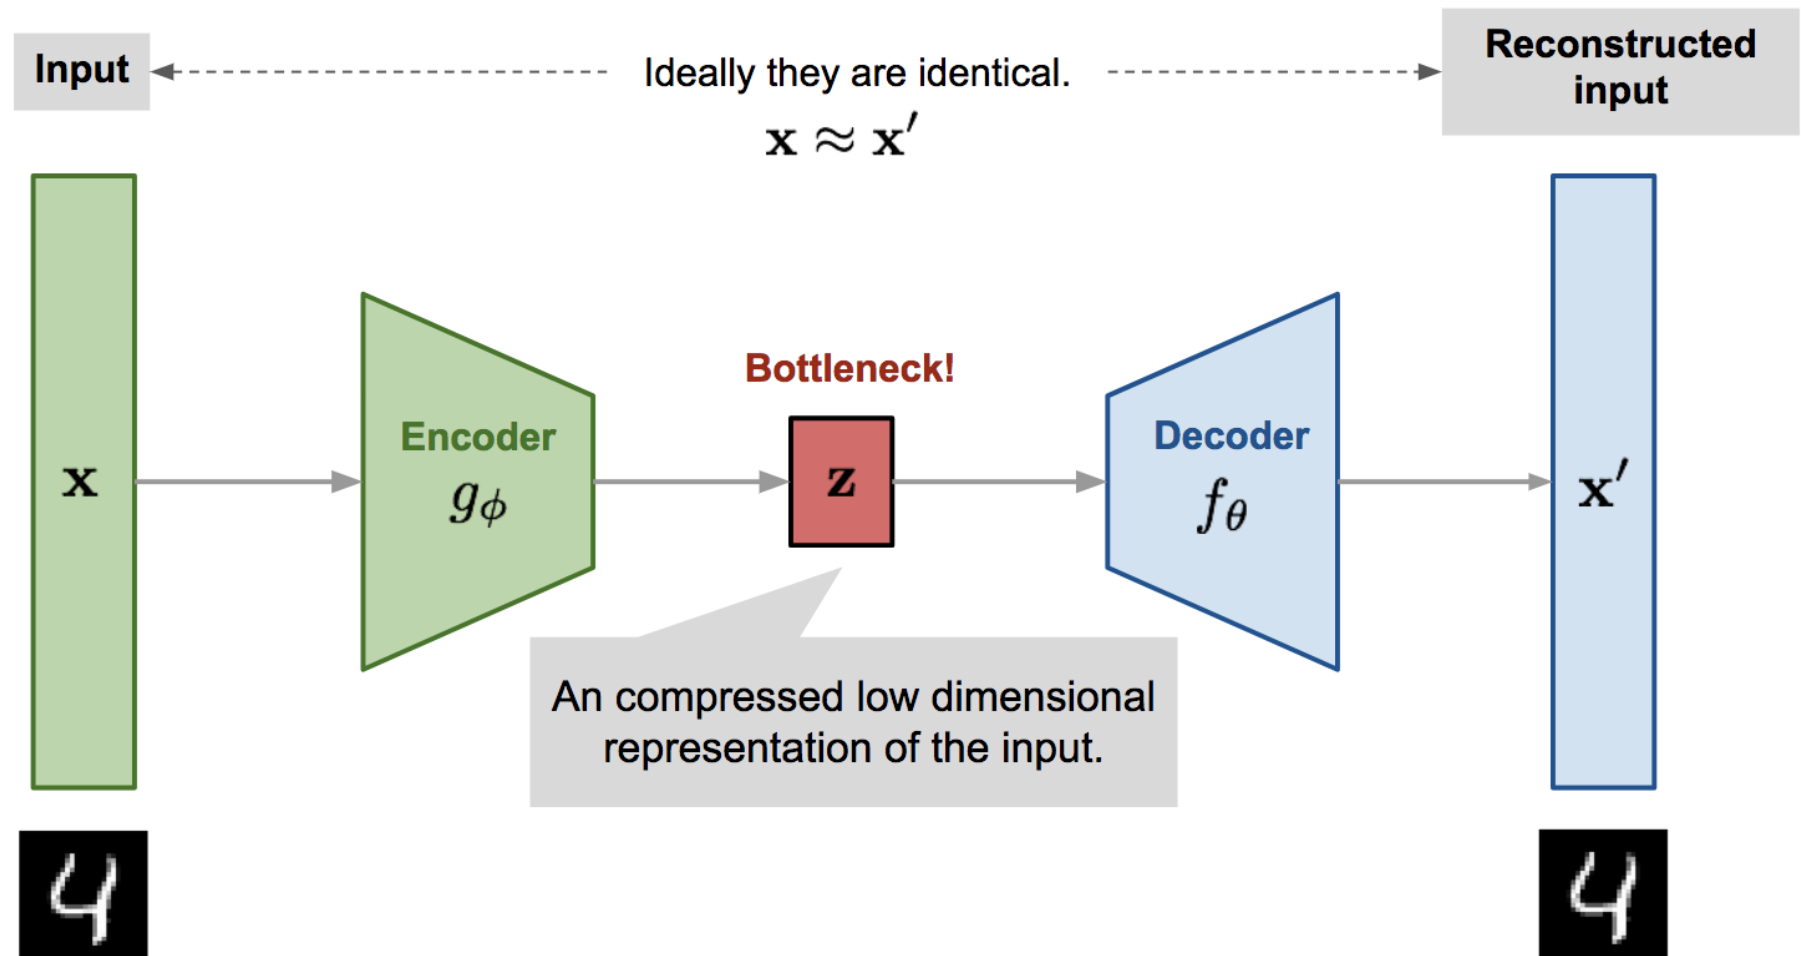
Recurrent Neural Networks (RNN)

- Sequence models, e.g., speech recognition, sentiment classification, ... etc.
- Inputs & outputs can have different lengths within the same dataset

Recurrent Neural Networks (RNN)



Autoencoder Network



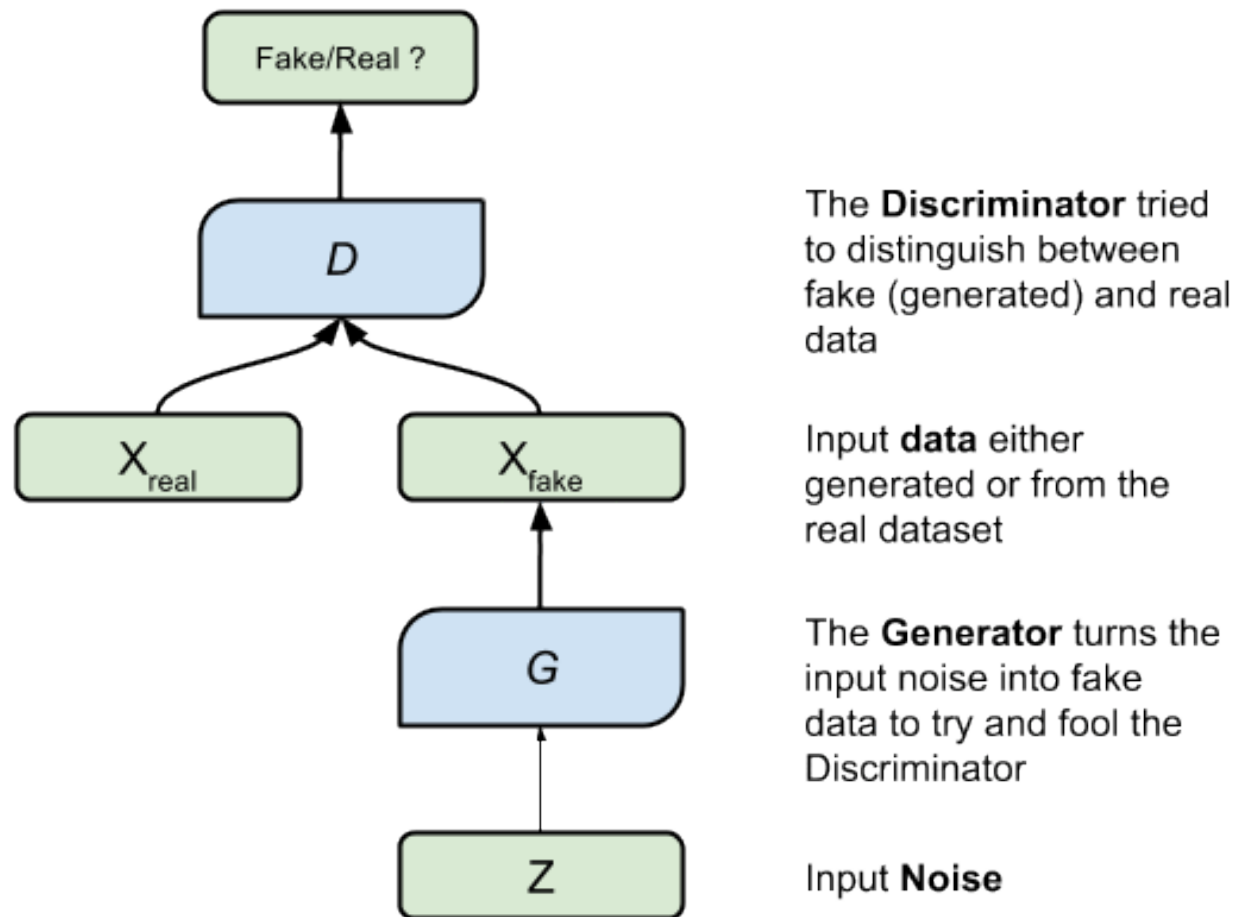
Source: Lilian Weng's Github blog

Autoencoder Network


- Unsupervised network
- Gives embedding
 - Better embeddings using supervised

Generative Adversarial Network (GAN)

- Create a generative model of artificial data



Siamese Networks




The **Distance Function** decides if the output vectors are close enough to be similar

The **Neural Network** transforms the input into a properties vector

Input Data (image, text, features...)

Source: Guy Ernest, AWS Amazon Blogs

Siamese Networks



The **Distance Function** decides if the output vectors are close enough to be similar

The **Neural Network** transforms the input into a properties vector

Input Data (image, text, features...)


Source: Guy Ernest, AWS Amazon Blogs

Deep Learning


- Subset of machine learning
- Multi-layered neural networks
- Raw data, i.e., end-to-end solution
- Requires big data & high computational power

Machine Learning vs Deep Learning


Machine Learning



Deep Learning

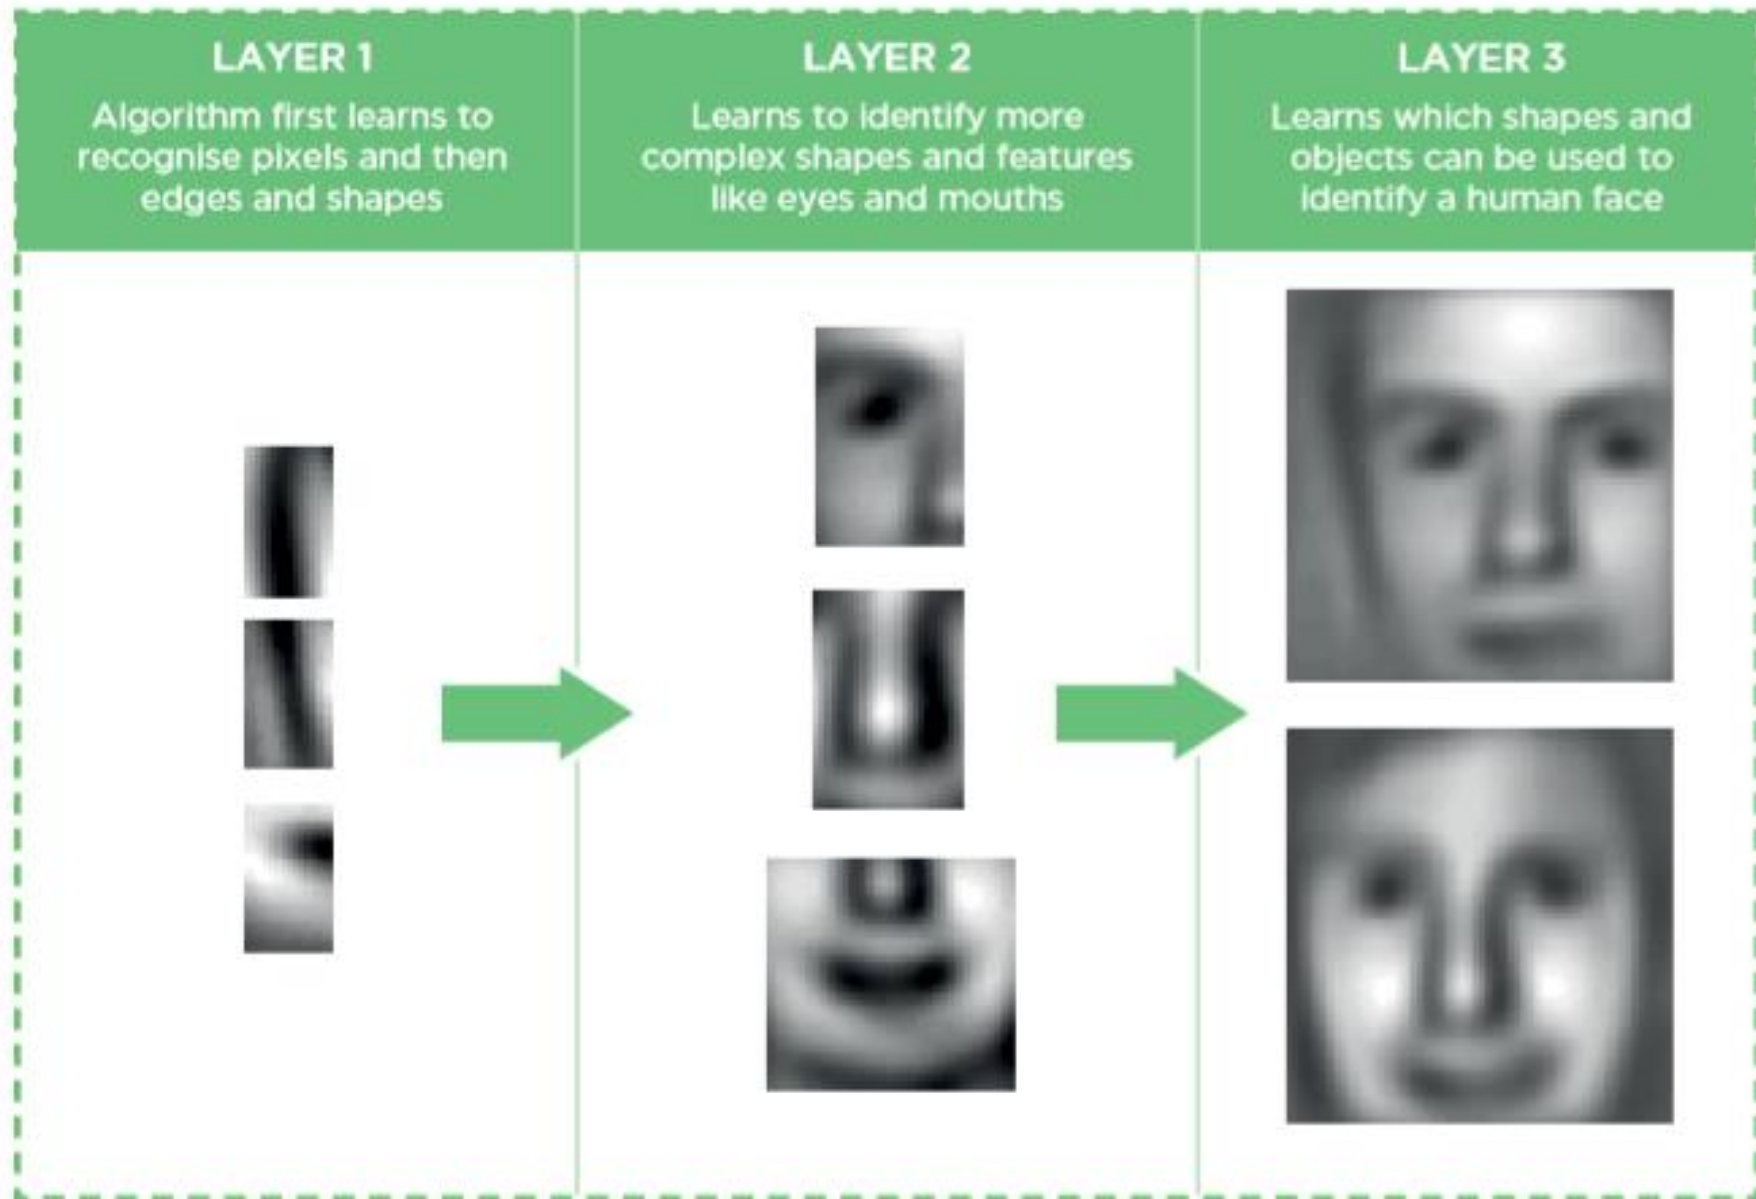


Machine Learning vs Deep Learning



Source: Hannes Schulz and Sven Behnke, 2012


Machine Learning vs Deep Learning




When to use Deep Learning?

- Big amount of **data** **expensive!**
- Availability of high computational power **expensive!**
- Lack of domain understanding
- Complex problems (vision, NLP, speech recognition)

Scalability with Data Amount



Scalability with Data Amount



Andrew Ng

Potentials of AI

"If a typical person can do a mental task with less than one second of thought, we can probably automate it using AI either now or in the near future."

— Andrew Ng

Currently, there are some limitations!

Limitations of Deep Learning

- Lots of achievements in vision field
- Not a magic tool!
 - Lack of adaptability and generality compared to human-vision system
 - Not able to build general-intelligent machine

Limitations of Deep Learning




Source: Gartner Hype Cycle for AI, 2019

Limitations of Deep Learning

- Why cannot fit all real-world scenarios?



Source: Boston Dynamics



Source: Google

Limitations of Deep Learning

- Large amount of labeled data
 - Impressive achievements correspond to supervised learning
 - Expensive!
 - Sometimes experts & special equipment are needed

Limitations of Deep Learning

- Datasets may be biased
 - Deep Networks become biased against rare patterns
 - Serious consequences in some real-world applications (e.g., medical, automotive, ... etc.)
 - Researchers should consider synthetic generation of data to mitigate the unbalanced representation of data

Limitations of Deep Learning

- Datasets may be biased



- Classification may be sensitive to viewpoint
 - if one of the viewpoints is under-represented


Limitations of Deep Learning

- Sensitive to standard adversarial attacks
 - Datasets are finite and just represent a fraction of all possible images



king penguin

+



adversarial perturbation

=


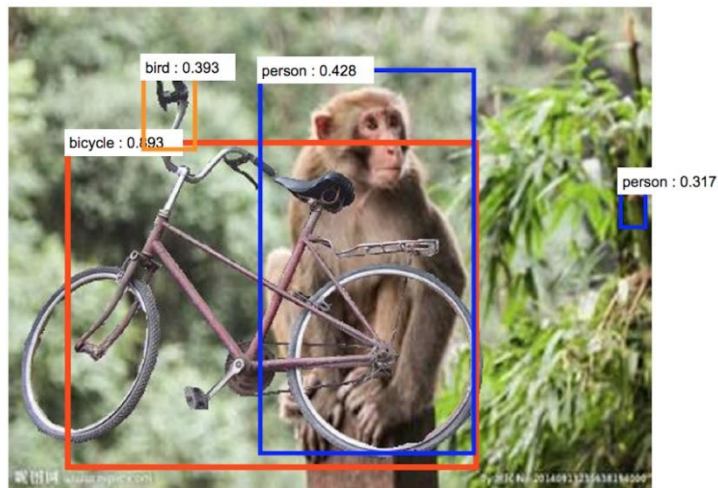
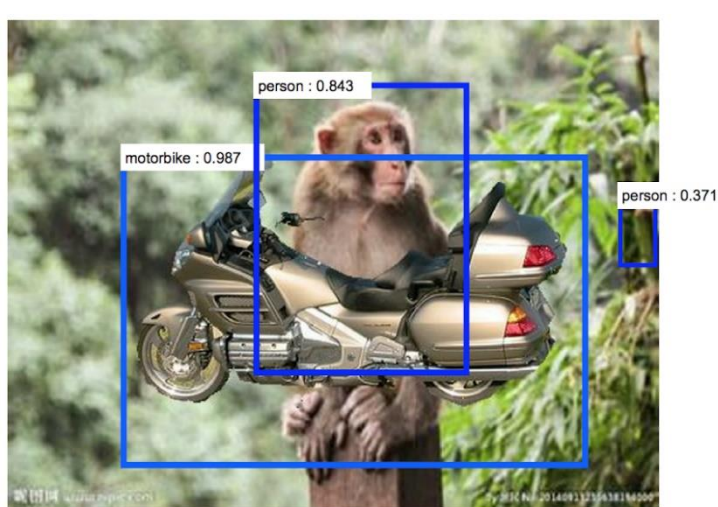


chihuahua

- Add extra training, i.e., “adversarial training”

Limitations of Deep Learning

- Over-sensitive to changes in context
 - Limited number of contexts in dataset, i.e., monkey in jungle
 - Combinatorial Explosion!



Limitations of Deep Learning

- Combinatorial Explosion
 - Real world images are combinatorial large
 - Application dependent (e.g., medical imaging is an exception)
 - Considering compositionality may be a potential solution
 - Testing is challenging (consider worst case scenarios)

Limitations of Deep Learning

- Visual understanding is tricky
 - Mirrors
 - Sparse Information
 - Physics
 - Humor
- Unintended results from fitness functions

Machine Learning Model Selection

1. Categorize the problem:

- Input: supervised, unsupervised, ... etc.
- Output: numerical → regression, class → classification, set of input groups → clustering

Machine Learning Model Selection

2. Understand your data:

a) Analyze the data:

- Descriptive statistics
- Data visualization

b) Process the data:

- Pre-processing, cleansing, ... etc.

c) Feature Engineering

Machine Learning Model Selection

3. Determine the possible algorithms:

- Based on categorization & data understanding
- May have a look at the literature
- Determine: desired accuracy, interpretability, scalability, complexity, training & testing time, runtime, ... etc.

Machine Learning Model Selection

4. Implement Machine Learning Algorithms:


- Setup a pipeline
- Compare algorithms
- Select an evaluation criteria

5. Tune hyperparameters

Time Series Prediction

- Time series contains:

- Trend
- Seasonality




- De-seasonalization:

- Remove the seasonal periodicities

TS → deseasonalize → predict →
return back
seasonal
components

How to deseasonalize?

- Removing the seasonal periodicities
- Usually seasonal cycle length is 12 months




How to deseasonalize?

- Obtain average of TS values over this window

$$a(year) = \frac{1}{12} \sum_{window} x(t)$$

- Normalization step: $Z(i) = \frac{x(i)}{a(year)}$
- Seasonal average \equiv avg of $Z(i)$'s of the different years for month i


$$u(i) = \frac{\sum_j Z_j(i)}{\#years}$$



How to deseasonalize?

- Seasonal average \equiv avg of $Z(i)$'s of the different years for month i

$$u(i) = \frac{\sum_{j=1}^{\#years} Z_j(i)}{\#years}$$




Deseasonalization Step

- Divide time series value by the corresponding seasonal average

$$x_{deseasonal}(t) = \frac{x(t)}{u(month(t))}$$

- After that focus on predicting the trend



Recover Seasonality

- After trend prediction, seasonality can be recovered via multiplication by the corresponding seasonal average

Acknowledgment

- These slides have been created relying on lecture notes of Prof. Dr. Amir Atiya